

Online Methods

Last Time

- Does value iteration always converge?
- Is the value function unique?


Guiding Questions

- What are the differences between *online* and *offline* solutions?
- Are there solution techniques that require computation time *independent* of the state space size?

Why Do We Need Something Else?

- Problems Policy and Value Iteration may struggle with?
 - Path planning across the country, or interplanetary
 - More realistic car dynamics (continuous states)
- Why are these problems hard?
 - State Space is massive (or infinite)

Curse of Dimensionality



1 dimension

e.g. $s = x \in S = \{1, 2, 3, 4, 5\}$

$$|S| = 5$$

2 dimensions

e.g. $s = (x, y) \in S = \{1, 2, 3, 4, 5\}^2$


$$|S| = 25$$

3 dimensions

e.g. $s = (x, y, x_h) \in S = \{1, 2, 3, 4, 5\}^3$

$$|S| = 125$$

(Discretize each dimension into 5 segments)




d dimensions, k segments $\rightarrow |S| = k^d$

Offline vs Online Solutions


Offline

- Before Execution: find V^*/Q^*
- During Execution: $\pi^*(s) = \text{argmax } Q^*(s, a)$



Online

- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)



- Why?
- Online methods are insensitive to the size of S !

```

struct MDP
    γ # discount factor
    S # state space
    A # action space
    T # transition function
    R # reward function
    TR # sample transition and reward
end

```

Rollout Lookahead

```

struct RolloutLookahead
    P # problem
    π # rollout policy
    d # depth
end

randstep(P::MDP, s, a) = P.TR(s, a)

function rollout(P, s, π, d)
    ret = 0.0
    for t in 1:d
        a = π(s)
        s, r = randstep(P, s, a)
        ret += P.γ^(t-1) * r
    end
    return ret
end

function (π::RolloutLookahead)(s)
    U(s) = rollout(π.P, s, π.π, π.d)
    return greedy(π.P, U, s).a
end

function greedy(P::MDP, U, s)
    u, a = findmax(a → lookahead(P, U, s, a), P.A)
    return (a=a, u=u)
end

function lookahead(P::MDP, U, s, a)
    S, T, R, γ = P.S, P.T, P.R, P.γ
    return R(s,a) + γ*sum(T(s,a,s')*U(s') for s' in S)

```


Forward Search

Algorithm 4.6 Forward search

```
1: function SELECTACTION( $s, d$ )
2:   if  $d = 0$ 
3:     return ( $\text{NIL}, 0$ )
4:    $(a^*, v^*) \leftarrow (\text{NIL}, -\infty)$ 
5:   for  $a \in A(s)$ 
6:      $v \leftarrow R(s, a)$ 
7:     for  $s' \in S(s, a)$ 
8:        $(a', v') \leftarrow \text{SELECTACTION}(s', d - 1)$ 
9:        $v \leftarrow v + \gamma T(s' \mid s, a)v'$ 
10:      if  $v > v^*$ 
11:         $(a^*, v^*) \leftarrow (a, v)$ 
12:   return ( $a^*, v^*$ )
```

$$(|S| \times |A|)^d$$

Forward Search depth



Sparse Sampling

Algorithm 4.8 Sparse sampling

```
1: function SELECTACTION( $s, d$ )
2:   if  $d = 0$ 
3:     return ( $\text{NIL}, 0$ )
4:    $(a^*, v^*) \leftarrow (\text{NIL}, -\infty)$ 
5:   for  $a \in A(s)$ 
6:      $v \leftarrow 0$ 
7:     for  $i \leftarrow 1$  to  $n$ 
8:        $(s', r) \sim G(s, a)$ 
9:        $(a', v') \leftarrow \text{SELECTACTION}(s', d - 1)$ 
10:       $v \leftarrow v + (r + \gamma v')/n$ 
11:      if  $v > v^*$ 
12:         $(a^*, v^*) \leftarrow (a, v)$ 
13:   return ( $a^*, v^*$ )
```

$$(n|A|)^d \quad |V^{\text{SS}}(s) - V^*(s)| \leq \epsilon \quad n, \epsilon, \text{ and } d \text{ related, but independent of } |S|$$

<https://www.cis.upenn.edu/~mkearns/papers/sparsesampling-journal.pdf>

Break

Draw the trees produced by the following algorithms for a problem with 2 actions and 3 states:

1. One-step lookahead with rollout
2. Forward search ($d=2$)
3. Sparse sampling ($d=2, n=2$)

Monte Carlo Tree Search (MCTS/UCT)

Keep track of:

$Q(s, a)$: Value estimate of that state and action combo

$N(s, a)$: Number of times we visit a state and action combo

$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \quad Q(s, a) + c \frac{N(s)^\beta}{\sqrt{N(s, a)}}$$


low $N(s, a)/N(s)$ = high bonus
start with $c = 2(\bar{V} - \underline{V})$, $\beta = 1/4$


Full story can be found in
<https://arxiv.org/pdf/1902.05213.pdf>

Monte Carlo Tree Search (MCTS/UCT)

```
function (π::MonteCarloTreeSearch)(s)
    for k in 1:π.m
        simulate!(π, s)
    end
    return argmax(a→π.Q[(s,a)], π.Π.ℳ)
end

function simulate!(π::MonteCarloTreeSearch, s, d=π.d)
    if d ≤ 0
        return π.U(s)
    end
    ℙ, N, Q, c = π.ℙ, π.N, π.Q, π.c
    ℳ, TR, γ = ℙ.ℳ, ℙ.TR, ℙ.γ
    if !haskey(N, (s, first(ℳ)))
        for a in ℳ
            N[(s,a)] = 0
            Q[(s,a)] = 0.0
        end
        return π.U(s)
    end
    a = explore(π, s)
    s', r = TR(s,a)
    q = r + γ*simulate!(π, s', d-1)
    N[(s,a)] += 1
    Q[(s,a)] += (q-Q[(s,a)])/N[(s,a)]
    return q
end
```





Using Online Methods in a Simulation

Algorithm: Rollout Simulation

Given: MDP (S, A, R, T, γ, b)

$s \leftarrow \text{sample}(b)$

$\hat{u} \leftarrow 0$

for t in $0 \dots T - 1$

$a \leftarrow \pi(s)$

$s', r \leftarrow G(s, a)$

$\hat{u} \leftarrow \hat{u} + \gamma^t r$

$s \leftarrow s'$

return \hat{u}

Guiding Questions

- What are the differences between online and offline solutions?
- Are there solution techniques that are *independent* of the state space size?

Forward Search Sparse Sampling

(FSSS)

Paper: <https://cdn.aaai.org/ojs/7689/7689-13-11219-1-2-20201228.pdf>

- Sparse Sampling, but only look at potentially valuable states

Things it keeps track of:

$Q(s, a)$: Estimate of the value for the state action pair

$U(s)$: Upper bound for value of state s

$L(s)$: Lower bound for value of state s

$U(s, a)$: Upper bound for value of state-action

$L(s, a)$: Lower bound for value of state-action

Forward Search Sparse Sampling

Algorithm 3 FSSS(s, d)

```
if  $d = 1$  (leaf) then
     $L^d(s, a) = U^d(s, a) = R(s, a), \forall a$ 
     $L^d(s) = U^d(s) = \max_a R(s, a)$ 
else if  $n_{sd} = 0$  then
    for each  $a \in A$  do
         $L^d(s, a) = V_{\min}$ 
         $U^d(s, a) = V_{\max}$ 
    for C times do
         $s' \sim T(s, a, \cdot)$ 
         $L^{d-1}(s') = V_{\min}$ 
         $U^{d-1}(s') = V_{\max}$ 
         $K^d(s, a) = K^d(s, a) \cup \{s'\}$ 
     $a^* = \operatorname{argmax}_a U^d(s, a)$ 
     $s^* = \max_{s' \in K^d(s, a^*)} (U^{d-1}(s') - L^{d-1}(s'))$ 
    FSSS( $s^*, d - 1$ )
     $n_{sd} = n_{sd} + 1$ 
     $L^d(s, a^*) = R(s, a^*) + \gamma \sum_{s' \in K^d(s, a^*)} L^{d-1}(s') / C$ 
     $U^d(s, a^*) = R(s, a^*) + \gamma \sum_{s' \in K^d(s, a^*)} U^{d-1}(s') / C$ 
     $L^d(s) = \max_a L^d(s, a)$ 
     $U^d(s) = \max_a U^d(s, a)$ 
```

If $L(s, a^*) \geq \max_{a \neq a^*} U(s, a)$ for best action ($a^* = \operatorname{argmax}_a U(s, a)$):
then, the node is closed because the best action is found.