

# Online Methods

**Last Time**

# Last Time

- Policy Iteration

# Last Time

- Policy Iteration
- Value Iteration


— Policy Evaluation  
Policy Improvement  
— Bellman's Operator

# Last Time

- Policy Iteration
- Value Iteration
- Does Value Iteration always converge?

# Last Time

- Policy Iteration
- Value Iteration
- Does Value Iteration always converge?
- Is the optimal value function unique?



# Guiding Questions

# Guiding Questions

- What are the differences between *online* and *offline* solutions?
- Are there solution techniques that require computation time *independent* of the state space size?



# Why Do We Need Something Else?


- Problems Policy and Value Iteration may struggle with?
- Why are these problems hard?

# Why Do We Need Something Else?

- Problems Policy and Value Iteration may struggle with?
  - Path planning across the country, or interplanetary
- Why are these problems hard?

# Why Do We Need Something Else?

- Problems Policy and Value Iteration may struggle with?
  - Path planning across the country, or interplanetary
  - More realistic car dynamics (continuous states)
- Why are these problems hard?



# Why Do We Need Something Else?

- Problems Policy and Value Iteration may struggle with?
  - Path planning across the country, or interplanetary
  - More realistic car dynamics (continuous states)
- Why are these problems hard?
  - State Space is massive (or infinite)

# Curse of Dimensionality

# Curse of Dimensionality

1 dimension, 5 segments

$$|\mathcal{S}| = 5$$




# Curse of Dimensionality

1 dimension, 5 segments

$$|\mathcal{S}| = 5$$

2 dimensions, 5 segments

$$|\mathcal{S}| = 25$$



# Curse of Dimensionality

1 dimension, 5 segments


$$|\mathcal{S}| = 5$$

2 dimensions, 5 segments

$$|\mathcal{S}| = 25$$

3 dimensions, 5 segments

$$|\mathcal{S}| = 125$$





# Curse of Dimensionality

1 dimension, 5 segments


$$|\mathcal{S}| = 5$$

2 dimensions, 5 segments

$$|\mathcal{S}| = 25$$

3 dimensions, 5 segments

$$|\mathcal{S}| = 125$$



$$n \text{ dimensions, } k \text{ segments} \rightarrow |\mathcal{S}| = k^n$$

# Offline vs Online Solutions

Offline

Online

# Offline vs Online Solutions

## Offline

- Before Execution: find  $V^*/Q^*$

## Online

# Offline vs Online Solutions

## Offline


- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$

## Online

# Offline vs Online Solutions

## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$




## Online

# Offline vs Online Solutions

## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$




## Online

- Before Execution: <nothing>

# Offline vs Online Solutions

## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$




## Online

- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)

# Offline vs Online Solutions


## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$



## Online

- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)






# Offline vs Online Solutions


## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$



## Online

- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)




- Why?

# Offline vs Online Solutions


## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$




## Online

- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)




- Why?
- Online methods are insensitive to the size of  $S$  !

# One Step Lookahead



$$Q(s,a) = R(s,a) + \gamma E(U(s'))$$



```
randstep( $\mathcal{P}::\text{MDP}$ ,  $s$ ,  $a$ ) =  $\mathcal{P}.\text{TR}(s, a)$ 
```

```
function rollout( $\mathcal{P}$ ,  $s$ ,  $\pi$ ,  $d$ )
```

```
    ret = 0.0
```

```
    for t in 1:d
```

```
        a =  $\pi(s)$ 
```

```
         $s, r$  = randstep( $\mathcal{P}$ ,  $s$ ,  $a$ )
```

```
        ret +=  $\mathcal{P}.\gamma^{(t-1)} * r$ 
```

```
    end
```

```
    return ret
```

```
end
```

```
function ( $\pi::\text{RolloutLookahead}$ )
```

```
     $U(s)$  = rollout( $\pi.\mathcal{P}$ ,  $s$ ,  $\pi.\pi$ ,  $\pi.d$ )
```

```
    return greedy( $\pi.\mathcal{P}$ ,  $U$ ,  $s$ ).a
```

```
end
```

```
function greedy( $\mathcal{P}::\text{MDP}$ ,  $U$ ,  $s$ )
```

```
     $u, a$  = findmax( $a \rightarrow \text{lookahead}(\mathcal{P}, U, s, a), (\mathcal{P}.A)$ )
```

```
    return ( $a=a, u=u$ )
```

```
end
```

```
function lookahead( $\mathcal{P}::\text{MDP}$ ,  $U$ ,  $s$ ,  $a$ )
```

```
     $S, T, R, \gamma$  =  $\mathcal{P}.S, \mathcal{P}.T, \mathcal{P}.R, \mathcal{P}.\gamma$ 
```


```
    return  $R(s,a) + \gamma * \text{sum}(T(s,a,s') * U(s'))$  for  $s'$  in  $S$ 
```

run a sim  
return  $\hat{u}$

$\hat{\pi}$

depth

# Forward Search




```
function forward_search( $\mathcal{P}$ ,  $s$ ,  $d$ ,  $U$ )
    if  $d \leq 0$ 
        return (a=nothing,  $u=U(s)$ )
    end
    best = (a=nothing,  $u=-\text{Inf}$ )
     $U'(s) = \text{forward\_search}(\mathcal{P}, s, d-1, U).u$ 
    for a in  $\mathcal{P}.A$ 
         $u = \text{lookahead}(\mathcal{P}, U', s, a)$ 
        if  $u > \text{best}.u$ 
            best = (a=a,  $u=u$ )
        end
    end
    return best
end
```

```
function lookahead( $\mathcal{P}::\text{MDP}$ ,  $U$ ,  $s$ ,  $a$ )
     $S, T, R, \gamma = \mathcal{P}.S, \mathcal{P}.T, \mathcal{P}.R, \mathcal{P}.\gamma$ 
    return  $R(s,a) + \gamma * \text{sum}(T(s,a,s') * U(s'))$  for  $s'$  in  $S$ 
```


Size of tree  $(|S| \times |A|)^d$

# Forward Search depth

# Forward Search depth




# Forward Search depth



# Sparse Sampling

$$|S|=3$$



```
function sparse_sampling( $\mathcal{P}$ , s, d, m, U)
    if  $d \leq 0$ 
        return (a=nothing, u=U(s))
    end
    best = (a=nothing, u=-Inf)
    for a in  $\mathcal{P}.A$ 
        u = 0.0
        for i in 1:m
             $s', r$  = randstep( $\mathcal{P}$ , s, a)
             $a', u'$  = sparse_sampling( $\mathcal{P}$ ,  $s'$ , d-1, m, U)
             $u += (r + \mathcal{P}.\gamma * u') / m$ 
        end
        if  $u > \text{best}.u$ 
            best = (a=a, u=u)
        end
    end
    return best
end
```

size of  
tree  
 $(m|A|)^d$

$$|V^{\text{SS}}(s) - V^*(s)| \leq \epsilon$$




$m$ ,  $\epsilon$ , and  $d$  related, but independent of  $|S|$



# Break

Draw the trees produced by the following algorithms for a problem with 2 actions and 3 states:

1. One-step lookahead with rollout
2. Forward search ( $d=2$ )
3. Sparse sampling ( $d=2, m=2$ )



# Monte Carlo Tree Search (MCTS/UCT)

Keep track of:

$Q(s, a)$ : Value estimate of that  
state and action combo

$N(s, a)$ : Number of times we  
visit a state and action combo

# Monte Carlo Tree Search (MCTS/UCT)

Keep track of:

$Q(s, a)$ : Value estimate of that  
state and action combo

$N(s, a)$ : Number of times we  
visit a state and action combo

$$Q(s, a) + c\sqrt{\frac{\log N(s)}{N(s, a)}}$$

# Monte Carlo Tree Search (MCTS/UCT)

Keep track of:

$Q(s, a)$ : Value estimate of that  
state and action combo

$N(s, a)$ : Number of times we  
visit a state and action combo

$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}}$$

low  $N(s, a)/N(s)$  = high bonus

# Monte Carlo Tree Search (MCTS/UCT)

Keep track of:

$Q(s, a)$ : Value estimate of that  
state and action combo

$N(s, a)$ : Number of times we  
visit a state and action combo

$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}}$$

low  $N(s, a)/N(s)$  = high bonus


start with  $c = 2(\bar{V} - \underline{V})$ ,  $\beta = 1/4$

# Monte Carlo Tree Search (MCTS/UCT)

Keep track of:


$Q(s, a)$ : Value estimate of that state and action combo

$N(s, a)$ : Number of times we visit a state and action combo



$$N(s) = \sum_a N(s,a)$$

$$\hat{Q} = \underbrace{Q(s,a)}_{\text{Value}} + \underbrace{c \sqrt{\frac{\log N(s)}{N(s,a)}}}_{\text{exploration bonus}} \quad \hat{Q} = Q(s,a) + c \frac{N(s)^\beta}{\sqrt{N(s,a)}}$$



low  $N(s,a)/N(s)$  = high bonus  
start with  $c = 2(\bar{V} - \underline{V})$ ,  $\beta = 1/4$

Full story can be found in  
<https://arxiv.org/pdf/1902.05213.pdf>

# Monte Carlo Tree Search (MCTS/UCT)

```
function (π::MonteCarloTreeSearch)(s)
    for k in 1:π.m
        simulate!(π, s)
    end
    return argmax(a → π.Q[(s,a)], π.P.A)
end
```

*k iterations*

$k$  iterations

each trip down  
and back up is one  
iteration

At end choose action node with highest  $Q(s, a)$

```

function simulate!( $\pi$ ::MonteCarloTreeSearch, s, d= $\pi$ .d)
    if d  $\leq$  0
        return  $\pi$ .U(s)
    end
     $\mathcal{P}$ , N, Q, c =  $\pi$ . $\mathcal{P}$ ,  $\pi$ .N,  $\pi$ .Q,  $\pi$ .c
     $\mathcal{A}$ , TR,  $\gamma$  =  $\mathcal{P}$ . $\mathcal{A}$ ,  $\mathcal{P}$ .TR,  $\mathcal{P}$ . $\gamma$ 
    if !haskey(N, (s, first( $\mathcal{A}$ )))
        for a in  $\mathcal{A}$ 
            N[(s,a)] = 0
            Q[(s,a)] = 0.0
        end
    end
    return  $\pi$ .U(s)
end

a = explore( $\pi$ , s)
s', r = TR(s,a)
q = r +  $\gamma$ *simulate!( $\pi$ , s', d-1)
N[(s,a)] += 1
Q[(s,a)] += (q-Q[(s,a)])/N[(s,a)]
return q
end

```

Handwritten annotations on the right side of the code:

- 2. Expansion (red bracket next to the initialization of N and Q)
- 3. Rollout (green bracket next to the recursive call to simulate!)
- 1. Search (blue bracket next to the explore function call)
- Backup (orange bracket next to the update of N and Q)

Handwritten note in the middle:

$-Q(s,a) + \text{exploration bonus}$


## 2. Expansion

### 3. Rollout/Value Estimate

## 1. Search

## Back up


each  
step



# Monte Carlo Tree Search (MCTS/UCT)

```
function ( $\pi$ ::MonteCarloTreeSearch)(s)
  for k in 1: $\pi$ .m
    simulate!( $\pi$ , s)
  end
  return argmax( $a \rightarrow \pi.Q[(s,a)]$ ,  $\pi.P.A$ )
end
```

```
function simulate!( $\pi$ ::MonteCarloTreeSearch, s, d= $\pi$ .d)
  if d  $\leq$  0
    return  $\pi.U(s)$ 
  end
   $P, N, Q, c = \pi.P, \pi.N, \pi.Q, \pi.c$ 
   $A, TR, \gamma = P.A, P.TR, P.\gamma$ 
  if !haskey( $N$ , (s, first( $A$ )))
    for a in  $A$ 
       $N[(s,a)] = 0$ 
       $Q[(s,a)] = 0.0$ 
    end
    return  $\pi.U(s)$ 
  end
  a = explore( $\pi$ , s)
   $s', r = TR(s,a)$ 
  q = r +  $\gamma * \text{simulate}!(\pi, s', d-1)$ 
   $N[(s,a)] += 1$ 
   $Q[(s,a)] += (q - Q[(s,a)]) / N[(s,a)]$ 
  return q
end
```







# Monte Carlo Tree Search (MCTS/UCT)

```
function ( $\pi$ ::MonteCarloTreeSearch)(s)
  for k in 1: $\pi$ .m
    simulate!( $\pi$ , s)
  end
  return argmax( $a \rightarrow \pi.Q[(s,a)]$ ,  $\pi.P.A$ )
end
```

```
function simulate!( $\pi$ ::MonteCarloTreeSearch, s, d= $\pi$ .d)
  if d  $\leq$  0
    return  $\pi.U(s)$ 
  end
   $P, N, Q, c = \pi.P, \pi.N, \pi.Q, \pi.c$ 
   $A, TR, \gamma = P.A, P.TR, P.\gamma$ 
  if !haskey(N, (s, first( $A$ )))
    for a in  $A$ 
       $N[(s,a)] = 0$ 
       $Q[(s,a)] = 0.0$ 
    end
    return  $\pi.U(s)$ 
  end
  a = explore( $\pi$ , s)
   $s', r = TR(s,a)$ 
  q = r +  $\gamma$ *simulate!( $\pi$ ,  $s'$ , d-1)
   $N[(s,a)] += 1$ 
   $Q[(s,a)] += (q-Q[(s,a)])/N[(s,a)]$ 
  return q
end
```





# Using Online Methods in a Simulation

# Using Online Methods in a Simulation

## Algorithm: Rollout Simulation

Given: MDP  $(S, A, R, T, \gamma, b)$

$s \leftarrow \text{sample}(b)$

$\hat{u} \leftarrow 0$

for  $t$  in  $0 \dots T - 1$

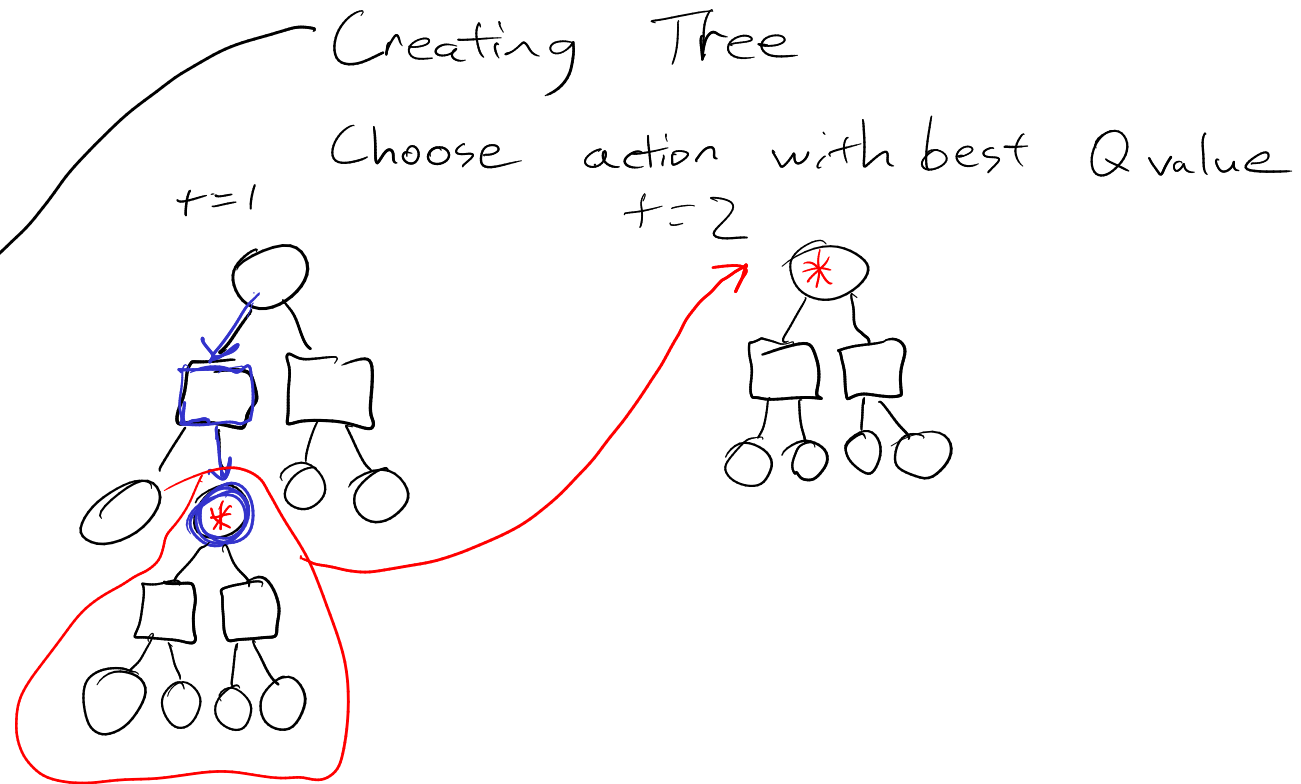
$a \leftarrow \boxed{\pi(s)}$

$s', r \leftarrow G(s, a)$

$\hat{u} \leftarrow \hat{u} + \gamma^t r$

$s \leftarrow s'$

return  $\hat{u}$



# Guiding Questions

# Guiding Questions

- What are the differences between online and offline solutions?
- Are there solution techniques that are *independent* of the state space size?

# Forward Search Sparse Sampling

(FSSS)

Paper: <https://cdn.aaai.org/ojs/7689/7689-13-11219-1-2-20201228.pdf>

- Sparse Sampling, but only look at potentially valuable states

# Forward Search Sparse Sampling

(FSSS)

Paper: <https://cdn.aaai.org/ojs/7689/7689-13-11219-1-2-20201228.pdf>

- Sparse Sampling, but only look at potentially valuable states

Things it keeps track of:

$Q(s, a)$ : Estimate of the value for the  
state action pair

$U(s)$ : Upper bound for value of state  $s$

$L(s)$ : Lower bound for value of state  $s$

$U(s, a)$ : Upper bound for value of state-  
action

$L(s, a)$ : Lower bound for value of state-  
action



# Forward Search Sparse Sampling

---

**Algorithm 3** FSSS( $s, d$ )

---

**if**  $d = 1$  (leaf) **then**

$$L^d(s, a) = U^d(s, a) = R(s, a), \forall a$$

$$L^d(s) = U^d(s) = \max_a R(s, a)$$

**else if**  $n_{sd} = 0$  **then**

**for each**  $a \in A$  **do**

$$L^d(s, a) = V_{\min}$$

$$U^d(s, a) = V_{\max}$$

**for** C times **do**

$$s' \sim T(s, a, \cdot)$$

$$L^{d-1}(s') = V_{\min}$$

$$U^{d-1}(s') = V_{\max}$$

$$K^d(s, a) = K^d(s, a) \cup \{s'\}$$

$$a^* = \operatorname{argmax}_a U^d(s, a)$$

$$s^* = \max_{s' \in K^d(s, a^*)} (U^{d-1}(s') - L^{d-1}(s'))$$

FSSS( $s^*, d - 1$ )

$$n_{sd} = n_{sd} + 1$$

$$L^d(s, a^*) = R(s, a^*) + \gamma \sum_{s' \in K^d(s, a^*)} L^{d-1}(s') / C$$

$$U^d(s, a^*) = R(s, a^*) + \gamma \sum_{s' \in K^d(s, a^*)} U^{d-1}(s') / C$$

$$L^d(s) = \max_a L^d(s, a)$$

$$U^d(s) = \max_a U^d(s, a)$$

---

# Forward Search Sparse Sampling

---

**Algorithm 3** FSSS( $s, d$ )

---

**if**  $d = 1$  (leaf) **then**

$$L^d(s, a) = U^d(s, a) = R(s, a), \forall a$$

$$L^d(s) = U^d(s) = \max_a R(s, a)$$

**else if**  $n_{sd} = 0$  **then**

**for each**  $a \in A$  **do**

$$L^d(s, a) = V_{\min}$$

$$U^d(s, a) = V_{\max}$$

**for**  $C$  times **do**

$$s' \sim T(s, a, \cdot)$$

$$L^{d-1}(s') = V_{\min}$$

$$U^{d-1}(s') = V_{\max}$$

$$K^d(s, a) = K^d(s, a) \cup \{s'\}$$

$$a^* = \operatorname{argmax}_a U^d(s, a)$$

$$s^* = \max_{s' \in K^d(s, a^*)} (U^{d-1}(s') - L^{d-1}(s'))$$

FSSS( $s^*, d - 1$ )

$$n_{sd} = n_{sd} + 1$$

$$L^d(s, a^*) = R(s, a^*) + \gamma \sum_{s' \in K^d(s, a^*)} L^{d-1}(s') / C$$

$$U^d(s, a^*) = R(s, a^*) + \gamma \sum_{s' \in K^d(s, a^*)} U^{d-1}(s') / C$$

$$L^d(s) = \max_a L^d(s, a)$$

$$U^d(s) = \max_a U^d(s, a)$$

---

If  $L(s, a^*) \geq \max_{a \neq a^*} U(s, a)$  for best action ( $a^* = \arg \max_a U(s, a)$ ):  
then, the node is closed because the best action is found.