

# Online Methods

# Last Time

# Last Time

- Does value iteration always converge?
- Is the value function unique?

# Guiding Questions

# Guiding Questions

- What are the differences between *online* and *offline* solutions?
- Are there solution techniques that require computation time *independent* of the state space size?

# Why Do We Need Something Else?

- Problems Policy and Value Iteration may struggle with?
- Why are these problems hard?

# Why Do We Need Something Else?

- Problems Policy and Value Iteration may struggle with?
  - Path planning across the country, or interplanetary
- Why are these problems hard?

# Why Do We Need Something Else?

- Problems Policy and Value Iteration may struggle with?
  - Path planning across the country, or interplanetary
  - More realistic car dynamics (continuous states)
- Why are these problems hard?

# Why Do We Need Something Else?

- Problems Policy and Value Iteration may struggle with?
  - Path planning across the country, or interplanetary
  - More realistic car dynamics (continuous states)
- Why are these problems hard?
  - State Space is massive (or infinite)

# Curse of Dimensionality



# Curse of Dimensionality

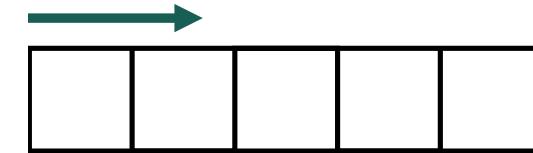


1 dimension

e.g.  $s = x \in S = \{1, 2, 3, 4, 5\}$

$$|S| = 5$$

(Discretize each dimension  
into 5 segments)



# Curse of Dimensionality



1 dimension

e.g.  $s = x \in S = \{1, 2, 3, 4, 5\}$

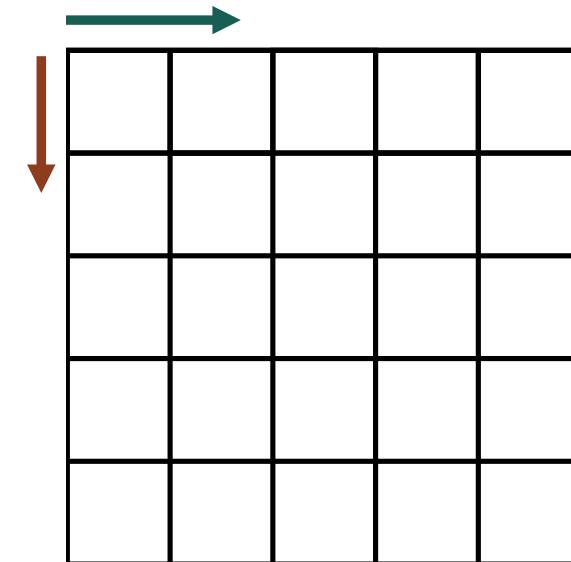
$$|S| = 5$$

2 dimensions

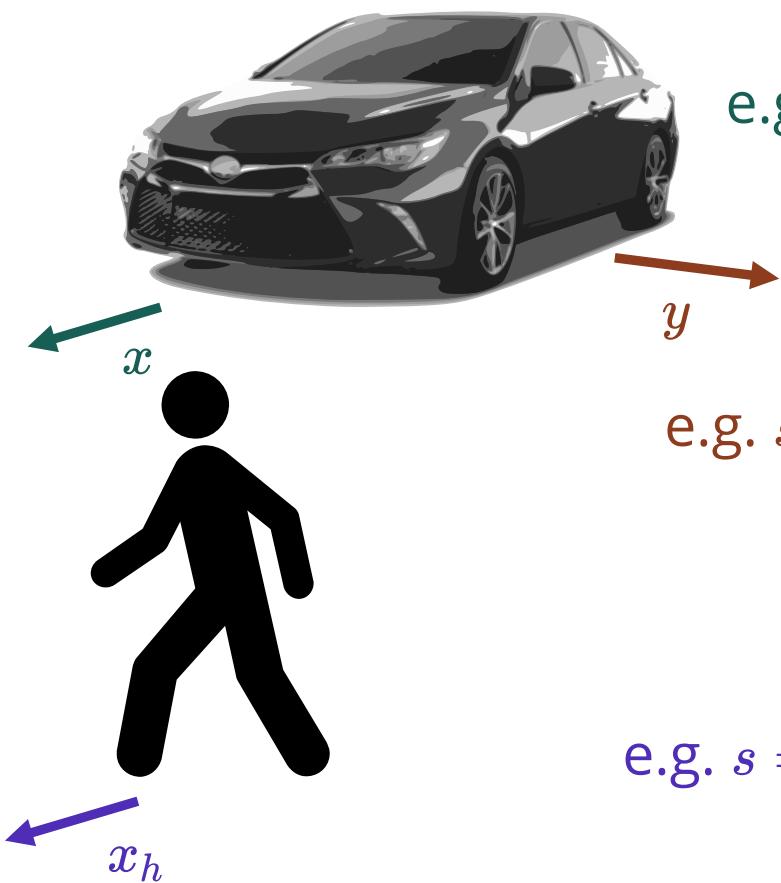
e.g.  $s = (x, y) \in S = \{1, 2, 3, 4, 5\}^2$

$$|S| = 25$$

(Discretize each dimension  
into 5 segments)



# Curse of Dimensionality



1 dimension

e.g.  $s = x \in S = \{1, 2, 3, 4, 5\}$

$$|S| = 5$$

2 dimensions

e.g.  $s = (x, y) \in S = \{1, 2, 3, 4, 5\}^2$

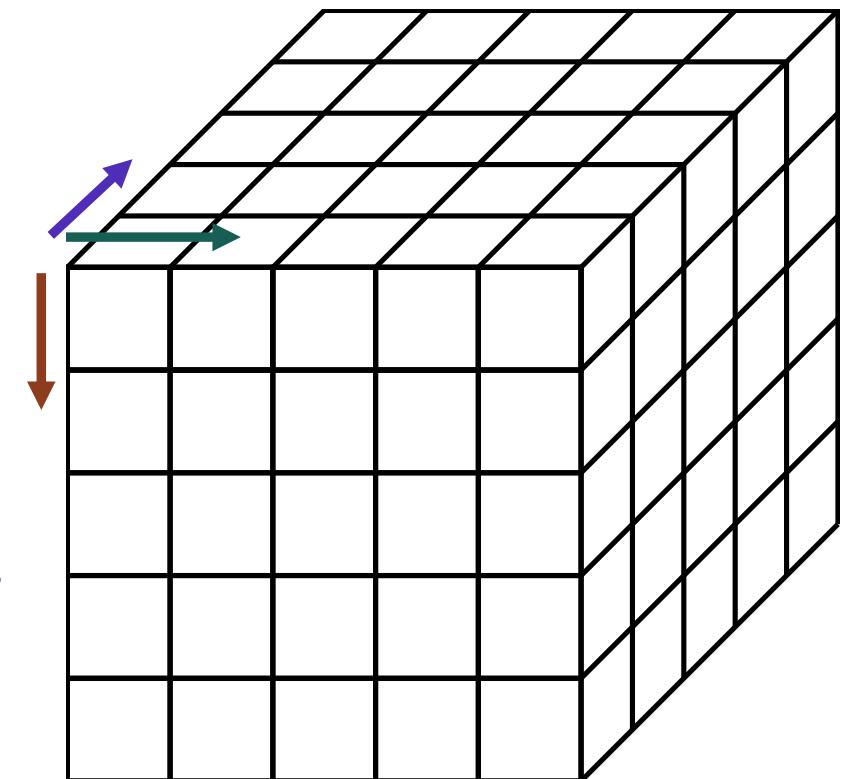
$$|S| = 25$$

3 dimensions

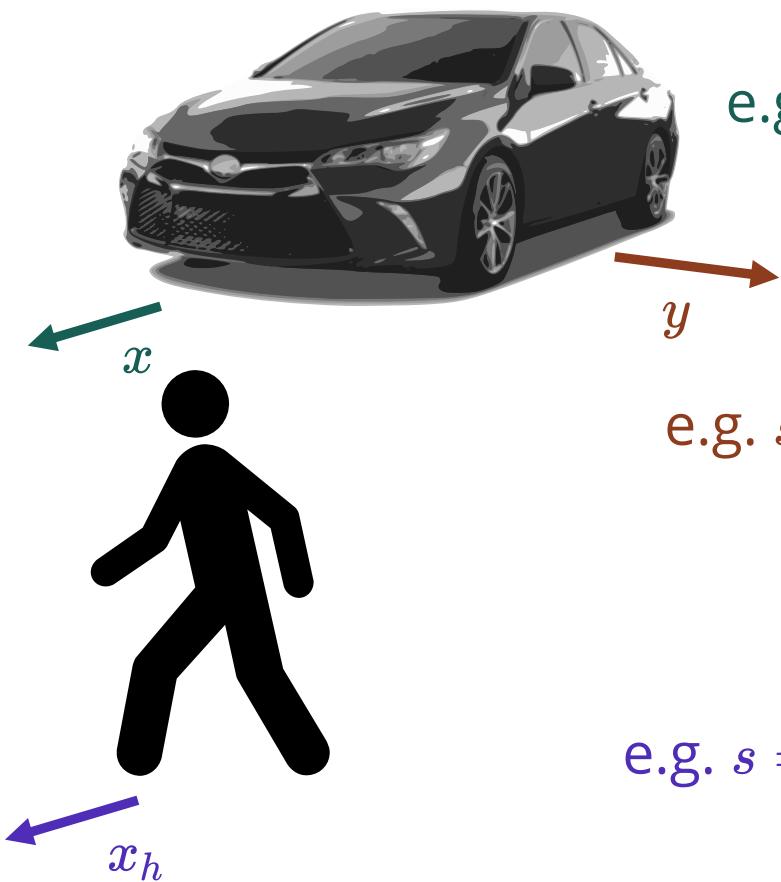
e.g.  $s = (x, y, x_h) \in S = \{1, 2, 3, 4, 5\}^3$

$$|S| = 125$$

(Discretize each dimension into 5 segments)



# Curse of Dimensionality



1 dimension

e.g.  $s = x \in S = \{1, 2, 3, 4, 5\}$

$$|S| = 5$$

2 dimensions

e.g.  $s = (x, y) \in S = \{1, 2, 3, 4, 5\}^2$

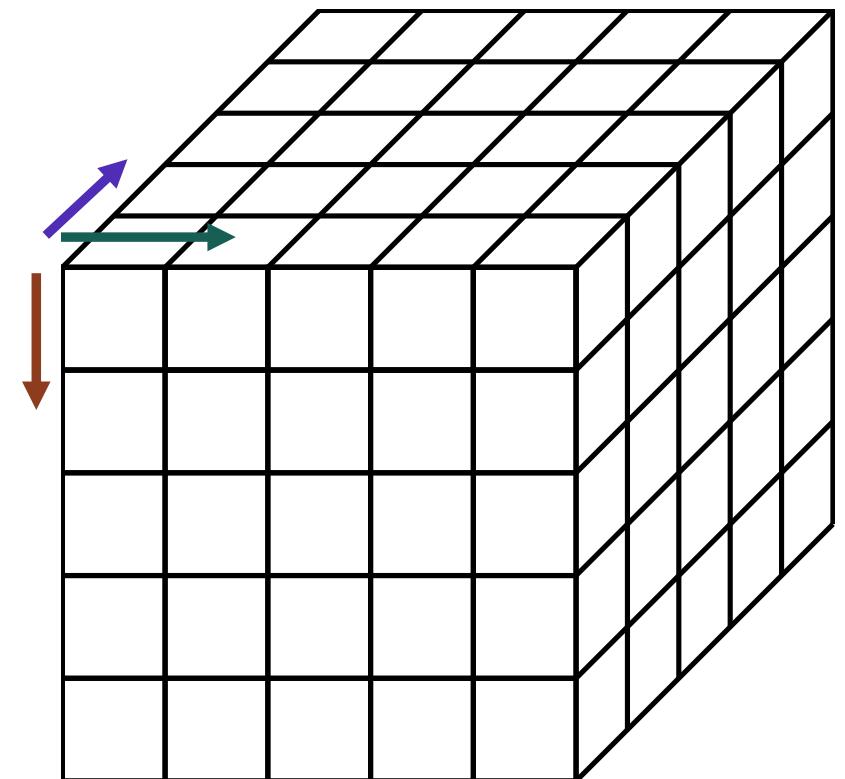
$$|S| = 25$$

3 dimensions

e.g.  $s = (x, y, x_h) \in S = \{1, 2, 3, 4, 5\}^3$

$$|S| = 125$$

(Discretize each dimension into 5 segments)



$d$  dimensions,  $k$  segments  $\rightarrow |S| = k^d$

# Offline vs Online Solutions

Offline

Online

# Offline vs Online Solutions

## Offline

- Before Execution: find  $V^*/Q^*$

## Online

# Offline vs Online Solutions

## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$

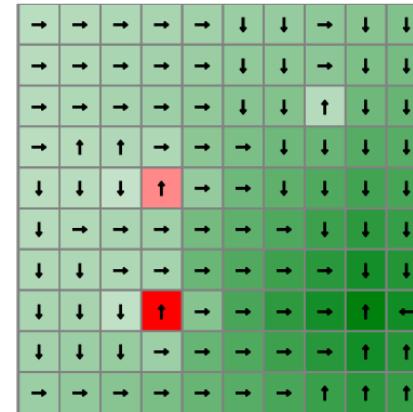
## Online

# Offline vs Online Solutions

## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$

## Online



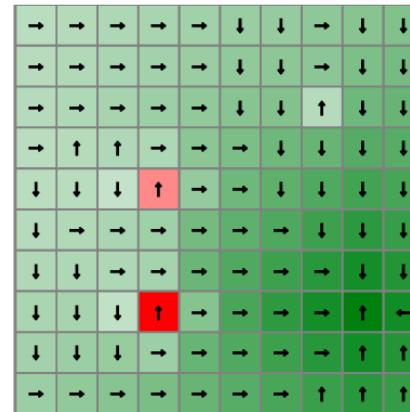
# Offline vs Online Solutions

## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$

## Online

- Before Execution: <nothing>



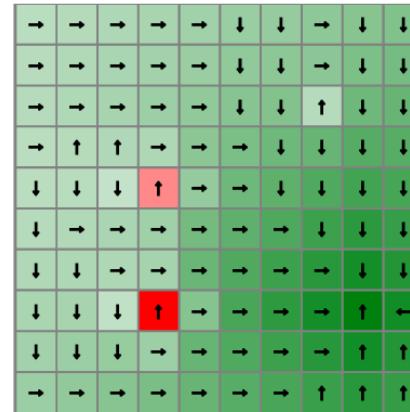
# Offline vs Online Solutions

## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$

## Online

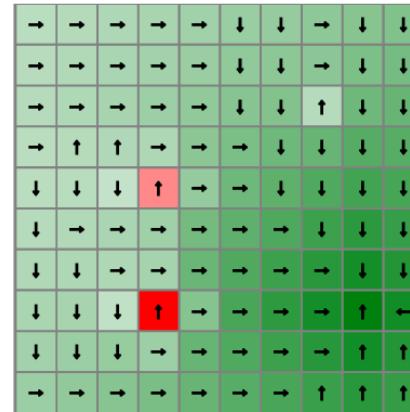
- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)



# Offline vs Online Solutions

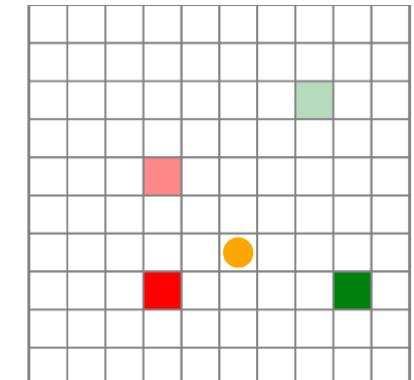
## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$



## Online

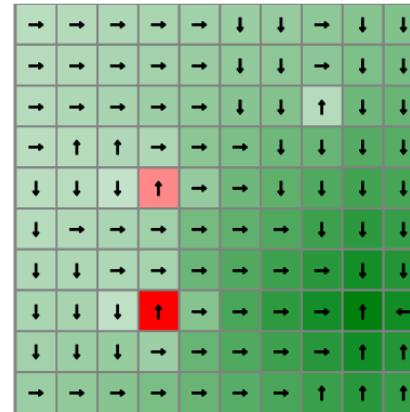
- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)



# Offline vs Online Solutions

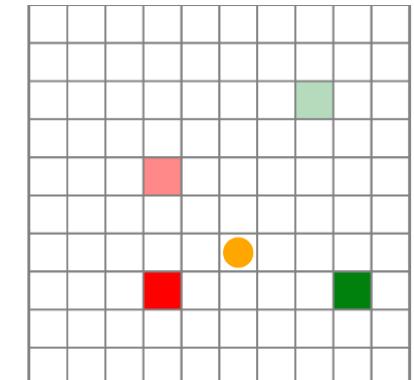
## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \text{argmax } Q^*(s, a)$



## Online

- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)

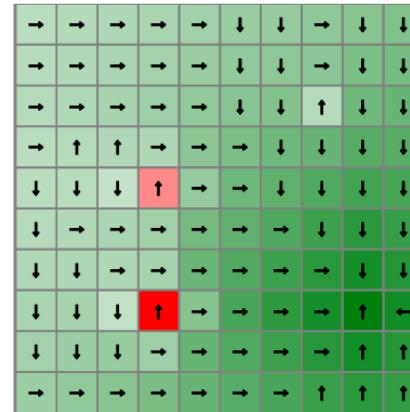


- Why?

# Offline vs Online Solutions

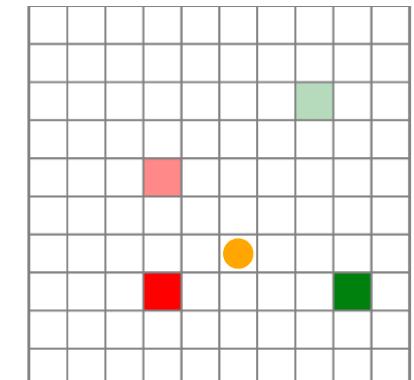
## Offline

- Before Execution: find  $V^*/Q^*$
- During Execution:  $\pi^*(s) = \text{argmax } Q^*(s, a)$



## Online

- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)



- Why?
- Online methods are insensitive to the size of  $S$  !

# One Step Lookahead

```
randstep( $\mathcal{P}$ ::MDP, s, a) =  $\mathcal{P}$ .TR(s, a)

function rollout( $\mathcal{P}$ , s,  $\pi$ , d)
    ret = 0.0
    for t in 1:d
        a =  $\pi$ (s)
        s, r = randstep( $\mathcal{P}$ , s, a)
        ret +=  $\mathcal{P}$ . $\gamma$ ^(t-1) * r
    end
    return ret
end

function ( $\pi$ ::RolloutLookahead)(s)
    U(s) = rollout( $\pi$ . $\mathcal{P}$ , s,  $\pi$ . $\pi$ ,  $\pi$ .d)
    return greedy( $\pi$ . $\mathcal{P}$ , U, s).a
end

function greedy( $\mathcal{P}$ ::MDP, U, s)
    u, a = findmax(a → lookahead( $\mathcal{P}$ , U, s, a),  $\mathcal{P}$ . $\mathcal{A}$ )
    return (a=a, u=u)
end

function lookahead( $\mathcal{P}$ ::MDP, U, s, a)
    S, T, R,  $\gamma$  =  $\mathcal{P}$ .S,  $\mathcal{P}$ .T,  $\mathcal{P}$ .R,  $\mathcal{P}$ . $\gamma$ 
    return R(s,a) +  $\gamma$ *sum(T(s,a,s')*U(s') for s' in S)
```

# Forward Search

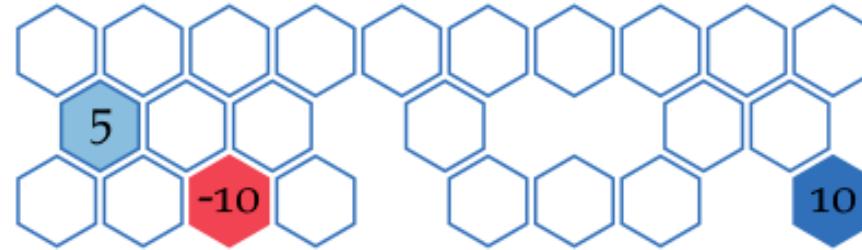
```
function forward_search(ℳ, s, d, U)
    if d ≤ 0
        return (a=nothing, u=U(s))
    end
    best = (a=nothing, u=-Inf)
    U'(s) = forward_search(ℳ, s, d-1, U).u
    for a in ℳ.ℳ
        u = lookahead(ℳ, U', s, a)
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end

function lookahead(ℳ::MDP, U, s, a)
    S, T, R, γ = ℳ.S, ℳ.T, ℳ.R, ℳ.γ
    return R(s,a) + γ*sum(T(s,a,s')*U(s') for s' in S)
```

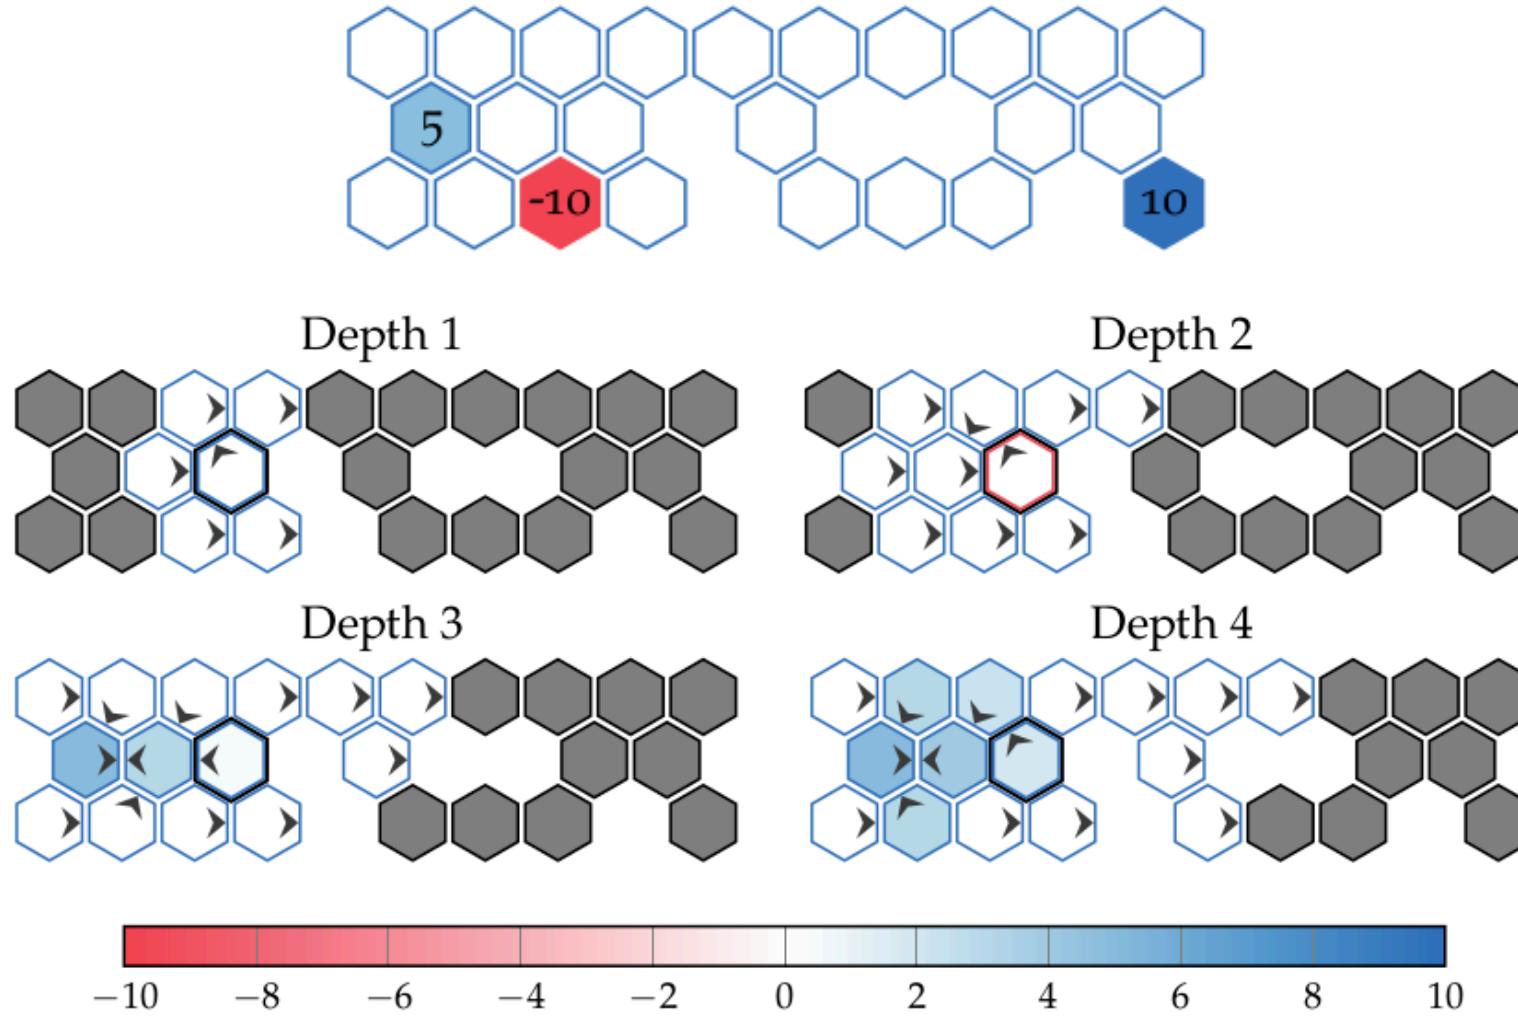
$$(|S| \times |A|)^d$$

# Forward Search depth

# Forward Search depth



# Forward Search depth



# Sparse Sampling

```
function sparse_sampling( $\mathcal{P}$ ,  $s$ ,  $d$ ,  $m$ ,  $U$ )
    if  $d \leq 0$ 
        return ( $a = \text{nothing}$ ,  $u = U(s)$ )
    end
    best = ( $a = \text{nothing}$ ,  $u = -\text{Inf}$ )
    for  $a$  in  $\mathcal{P}.\mathcal{A}$ 
         $u = 0.0$ 
        for  $i$  in  $1:m$ 
             $s'$ ,  $r = \text{randstep}(\mathcal{P}, s, a)$ 
             $a'$ ,  $u' = \text{sparse\_sampling}(\mathcal{P}, s', d-1, m, U)$ 
             $u += (r + \mathcal{P}.y * u') / m$ 
        end
        if  $u > \text{best}.u$ 
            best = ( $a = a$ ,  $u = u$ )
        end
    end
    return best
end
```

$$(m|A|)^d \quad |V^{\text{SS}}(s) - V^*(s)| \leq \epsilon \quad m, \epsilon, \text{ and } d \text{ related, but independent of } |S|$$

<https://www.cis.upenn.edu/~mkearns/papers/sparsesampling-journal.pdf>

# Break

Draw the trees produced by the following algorithms for a problem with 2 actions and 3 states:

1. One-step lookahead with rollout
2. Forward search ( $d=2$ )
3. Sparse sampling ( $d=2, m=2$ )

# Monte Carlo Tree Search (MCTS/UCT)

Keep track of:

$Q(s, a)$ : Value estimate of that state and action combo

$N(s, a)$ : Number of times we visit a state and action combo

# Monte Carlo Tree Search (MCTS/UCT)

Keep track of:

$Q(s, a)$ : Value estimate of that state and action combo

$N(s, a)$ : Number of times we visit a state and action combo

$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}}$$

# Monte Carlo Tree Search (MCTS/UCT)

Keep track of:

$Q(s, a)$ : Value estimate of that state and action combo

$N(s, a)$ : Number of times we visit a state and action combo

$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}}$$

low  $N(s, a)/N(s)$  = high bonus

# Monte Carlo Tree Search (MCTS/UCT)

Keep track of:

$Q(s, a)$ : Value estimate of that state and action combo

$N(s, a)$ : Number of times we visit a state and action combo

$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}}$$

low  $N(s, a)/N(s)$  = high bonus

start with  $c = 2(\bar{V} - \underline{V})$ ,  $\beta = 1/4$

# Monte Carlo Tree Search (MCTS/UCT)

Keep track of:

$Q(s, a)$ : Value estimate of that state and action combo

$N(s, a)$ : Number of times we visit a state and action combo

$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \quad Q(s, a) + c \frac{N(s)^\beta}{\sqrt{N(s, a)}}$$

low  $N(s, a)/N(s)$  = high bonus  
start with  $c = 2(\bar{V} - \underline{V})$ ,  $\beta = 1/4$

Full story can be found in  
<https://arxiv.org/pdf/1902.05213.pdf>

# Monte Carlo Tree Search (MCTS/UCT)

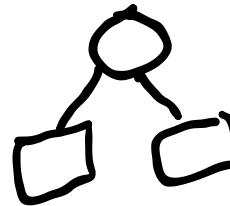
```
function (π::MonteCarloTreeSearch)(s)
    for k in 1:π.m
        simulate!(π, s)
    end
    return argmax(a→π.Q[(s,a)], π.Ρ.Α)
end

function simulate!(π::MonteCarloTreeSearch, s, d=π.d)
    if d ≤ 0
        return π.U(s)
    end
    Ρ, N, Q, c = π.Ρ, π.N, π.Q, π.c
    Α, TR, γ = Ρ.Α, Ρ.TR, Ρ.γ
    if !haskey(N, (s, first(Α)))
        for a in Α
            N[(s,a)] = 0
            Q[(s,a)] = 0.0
        end
        return π.U(s)
    end
    a = explore(π, s)
    s', r = TR(s,a)
    q = r + γ*simulate!(π, s', d-1)
    N[(s,a)] += 1
    Q[(s,a)] += (q-Q[(s,a)])/N[(s,a)]
    return q
end
```

# Monte Carlo Tree Search (MCTS/UCT)

```
function (π::MonteCarloTreeSearch)(s)
    for k in 1:π.m
        simulate!(π, s)
    end
    return argmax(a→π.Q[(s,a)], π.Ρ.Α)
end

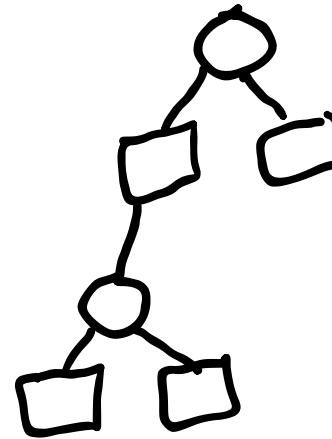
function simulate!(π::MonteCarloTreeSearch, s, d=π.d)
    if d ≤ 0
        return π.U(s)
    end
    Ρ, N, Q, c = π.Ρ, π.N, π.Q, π.c
    Α, TR, γ = Ρ.Α, Ρ.TR, Ρ.γ
    if !haskey(N, (s, first(Α)))
        for a in Α
            N[(s,a)] = 0
            Q[(s,a)] = 0.0
        end
        return π.U(s)
    end
    a = explore(π, s)
    s', r = TR(s,a)
    q = r + γ*simulate!(π, s', d-1)
    N[(s,a)] += 1
    Q[(s,a)] += (q-Q[(s,a)])/N[(s,a)]
    return q
end
```

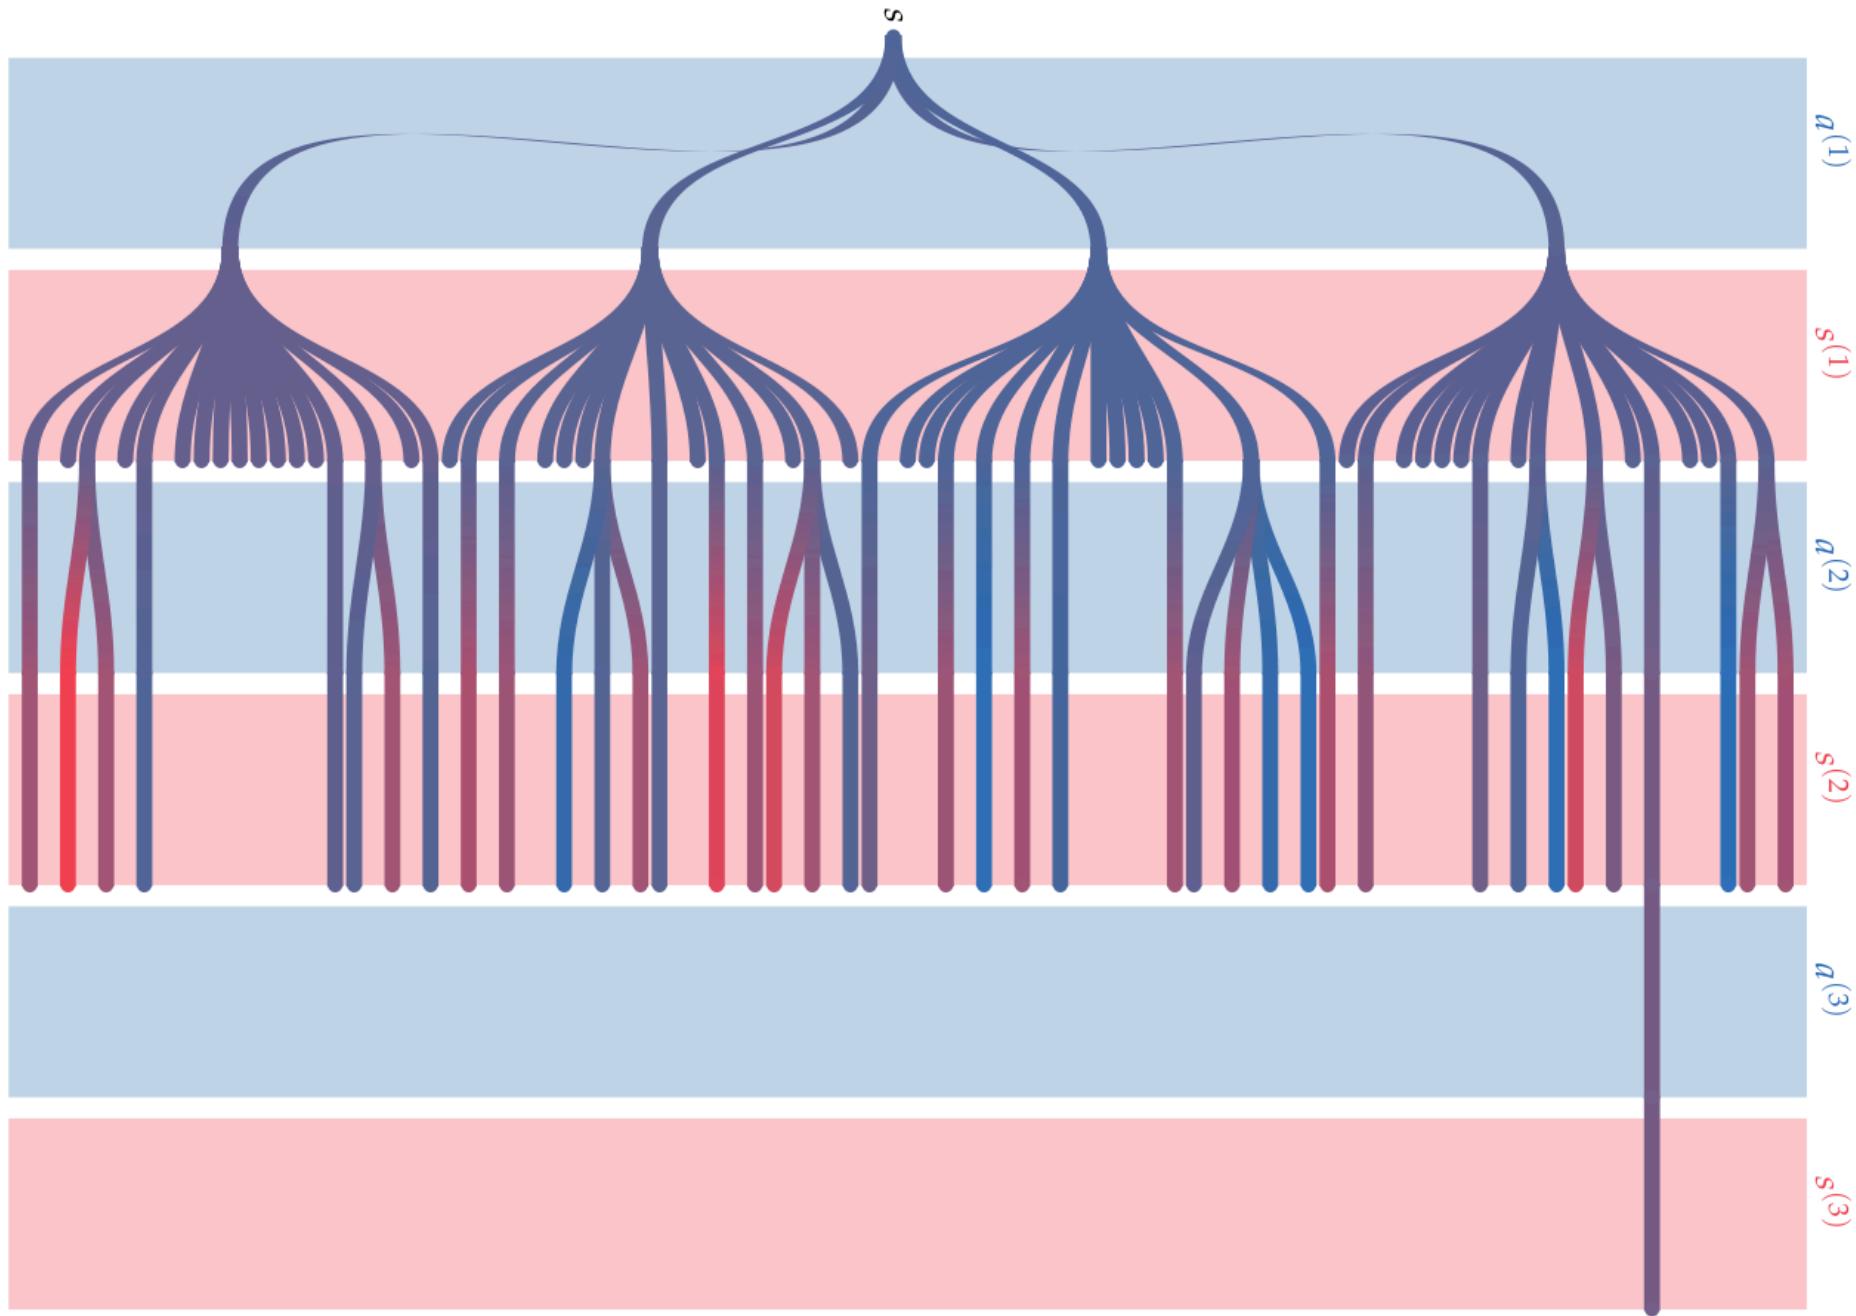


# Monte Carlo Tree Search (MCTS/UCT)

```
function (π::MonteCarloTreeSearch)(s)
    for k in 1:π.m
        simulate!(π, s)
    end
    return argmax(a→π.Q[(s,a)], π.Π.ℳ)
end

function simulate!(π::MonteCarloTreeSearch, s, d=π.d)
    if d ≤ 0
        return π.U(s)
    end
    ℙ, N, Q, c = π.ℙ, π.N, π.Q, π.c
    ℳ, TR, γ = ℙ.ℳ, ℙ.TR, ℙ.γ
    if !haskey(N, (s, first(ℳ)))
        for a in ℳ
            N[(s,a)] = 0
            Q[(s,a)] = 0.0
        end
        return π.U(s)
    end
    a = explore(π, s)
    s', r = TR(s,a)
    q = r + γ*simulate!(π, s', d-1)
    N[(s,a)] += 1
    Q[(s,a)] += (q-Q[(s,a)])/N[(s,a)]
    return q
end
```





# Using Online Methods in a Simulation

# Using Online Methods in a Simulation

## Algorithm: Rollout Simulation

Given: MDP  $(S, A, R, T, \gamma, b)$

$s \leftarrow \text{sample}(b)$

$\hat{u} \leftarrow 0$

for  $t$  in  $0 \dots T - 1$

$a \leftarrow \pi(s)$

$s', r \leftarrow G(s, a)$

$\hat{u} \leftarrow \hat{u} + \gamma^t r$

$s \leftarrow s'$

return  $\hat{u}$

# Guiding Questions

# Guiding Questions

- What are the differences between online and offline solutions?
- Are there solution techniques that are *independent* of the state space size?

# Forward Search Sparse Sampling

(FSSS)

Paper: <https://cdn.aaai.org/ojs/7689/7689-13-11219-1-2-20201228.pdf>

- Sparse Sampling, but only look at potentially valuable states

# Forward Search Sparse Sampling

(FSSS)

Paper: <https://cdn.aaai.org/ojs/7689/7689-13-11219-1-2-20201228.pdf>

- Sparse Sampling, but only look at potentially valuable states

Things it keeps track of:

$Q(s, a)$ : Estimate of the value for the state action pair

$U(s)$ : Upper bound for value of state s

$L(s)$ : Lower bound for value of state s

$U(s, a)$ : Upper bound for value of state-action

$L(s, a)$ : Lower bound for value of state-action

# Forward Search Sparse Sampling

---

**Algorithm 3** FSSS( $s, d$ )

---

```
if  $d = 1$  (leaf) then
     $L^d(s, a) = U^d(s, a) = R(s, a), \forall a$ 
     $L^d(s) = U^d(s) = \max_a R(s, a)$ 
else if  $n_{sd} = 0$  then
    for each  $a \in A$  do
         $L^d(s, a) = V_{\min}$ 
         $U^d(s, a) = V_{\max}$ 
    for C times do
         $s' \sim T(s, a, \cdot)$ 
         $L^{d-1}(s') = V_{\min}$ 
         $U^{d-1}(s') = V_{\max}$ 
         $K^d(s, a) = K^d(s, a) \cup \{s'\}$ 
     $a^* = \operatorname{argmax}_a U^d(s, a)$ 
     $s^* = \max_{s' \in K^d(s, a^*)} (U^{d-1}(s') - L^{d-1}(s'))$ 
    FSSS( $s^*, d - 1$ )
     $n_{sd} = n_{sd} + 1$ 
     $L^d(s, a^*) = R(s, a^*) + \gamma \sum_{s' \in K^d(s, a^*)} L^{d-1}(s') / C$ 
     $U^d(s, a^*) = R(s, a^*) + \gamma \sum_{s' \in K^d(s, a^*)} U^{d-1}(s') / C$ 
     $L^d(s) = \max_a L^d(s, a)$ 
     $U^d(s) = \max_a U^d(s, a)$ 
```

---

# Forward Search Sparse Sampling

---

**Algorithm 3** FSSS( $s, d$ )

---

```
if  $d = 1$  (leaf) then
     $L^d(s, a) = U^d(s, a) = R(s, a), \forall a$ 
     $L^d(s) = U^d(s) = \max_a R(s, a)$ 
else if  $n_{sd} = 0$  then
    for each  $a \in A$  do
         $L^d(s, a) = V_{\min}$ 
         $U^d(s, a) = V_{\max}$ 
    for C times do
         $s' \sim T(s, a, \cdot)$ 
         $L^{d-1}(s') = V_{\min}$ 
         $U^{d-1}(s') = V_{\max}$ 
         $K^d(s, a) = K^d(s, a) \cup \{s'\}$ 
     $a^* = \operatorname{argmax}_a U^d(s, a)$ 
     $s^* = \max_{s' \in K^d(s, a^*)} (U^{d-1}(s') - L^{d-1}(s'))$ 
    FSSS( $s^*, d - 1$ )
     $n_{sd} = n_{sd} + 1$ 
     $L^d(s, a^*) = R(s, a^*) + \gamma \sum_{s' \in K^d(s, a^*)} L^{d-1}(s') / C$ 
     $U^d(s, a^*) = R(s, a^*) + \gamma \sum_{s' \in K^d(s, a^*)} U^{d-1}(s') / C$ 
     $L^d(s) = \max_a L^d(s, a)$ 
     $U^d(s) = \max_a U^d(s, a)$ 
```

---

If  $L(s, a^*) \geq \max_{a \neq a^*} U(s, a)$  for best action ( $a^* = \operatorname{argmax}_a U(s, a)$ ):  
then, the node is closed because the best action is found.